

Rapport de Stage

BiTEM group - HES-SO Genève

Léandre Catogni

Bachelor en Mathématiques, Informatique et
Sciences Numériques

Université de Genève
Mai 2025

Remerciements

Je souhaite tout d'abord remercier Julien Knafou, grâce à qui j'ai pu trouver ce stage, pour l'aide qu'il m'a accordé dans mon apprentissage.

Je suis également particulièrement reconnaissant pour la confiance que m'a accordé Patrick Ruch, directeur du stage. TODO : Completer

Contents

1	Introduction	4
1.1	Context	4
1.2	Motivations	4
1.3	Objectives	4
2	BiTEM group presentation	5
2.1	Biodiversity Monitoring via Question Answering	5
3	State of the art	7
3.1	Tokenization	7
3.1.1	Byte Pair Encoding	9
3.1.2	WordPiece	10
3.2	Word Embeddings and Word2Vec	12
3.2.1	Term Frequency–Inverse Document Frequency (TF-IDF)	12
3.2.2	Goal	13
3.2.3	Word2Vec	13
3.2.4	Negative Sampling	14
3.2.5	GloVe	14
3.3	Neural Networks	15
3.4	Deep Contextual Models	16
3.4.1	Recurrent Neural Networks (RNNs)	16
3.4.2	Long Short-Term Memory (LSTM)	16
3.4.3	Attention for Neural Networks	17
3.4.4	Seq2seq	17
3.4.5	Transformers	18
3.4.6	BERT	18
3.4.7	BERT Variations	19
3.4.8	Heads	19
3.4.9	Fine-tuning	20
3.4.10	Optimization Algorithms	20
3.5	Evaluation	21
3.5.1	Data splitting	21
3.5.2	K-Fold Cross Validation	22
3.5.3	Metrics	23
3.5.4	Hyperparameter Optimization	26
3.5.5	Statistical Test	26

4 Data	30
4.1 Description	30
4.2 Pre-Processing	30
4.3 Repartition	31
5 Méthodes	33
5.1 Overview	33
5.2 Gestion des données	34
5.2.1 Data Splitting	34
5.2.2 Handling Imbalanced Data	34
5.3 Architecture	35
6 Résultats	36
7 Analyse et discussion	39
7.0.1 Ablation	39
7.0.2 Contamination	39
7.1 Limitations	40
8 Conclusion	41
9 Bibliography	42
10 Annexes	43
10.0.1 Workflow	44

Chapter 1

Introduction

1.1 Context

Over the past decade, the exponential growth of scientific publications has created significant challenges for researchers seeking to stay abreast of developments in their fields. Advances in natural language processing (NLP) have opened new avenues for distilling large volumes of textual data into concise summaries and extracting high-quality, relevant information [[citeulike:123456](#)]. Tools such as transformer-based models enable state-of-the-art performance in tasks including summarization, information retrieval, and question answering.

NLP systems now play a critical role in many applications that underpin modern society, including spam detection, grammar correction, and search engines. These technologies have become indispensable, allowing users to focus on substantive content and gain efficiency in research workflows [[jurafsky2009speech](#)].

In this internship, conducted from March to July 2025, I collaborated with the BiTeM group to develop AI-powered services for biodiversity monitoring via question answering.

1.2 Motivations

The volume of scientific literature published annually has increased by over 8% in the life sciences alone, reaching more than two million articles per year as of 2024 [[nih:stats2024](#)]. Researchers studying specialized domains, such as island ecosystems in biodiversity, often face difficulty locating precise, high-quality data among a noisy corpus. This motivates the development of automated classifiers to filter relevant publications and thus accelerate scientific discovery.

1.3 Objectives

My project is part of the Biodiversity Monitoring via Question Answering and aims at developing an automated transformers-powered classifier.

Chapter 2

BiTEM group presentation

The BiTeM group (Bibliomics and Text Mining Group), in which I did this internship, is part of the Information Science Department of the HES-SO/HEG Geneva. This research group's open-science projects mostly focus on clinical and biological data.

It is also affiliated with the text mining group of the Swiss Institute of Bioinformatics (SIB), which is currently working on the Swiss Institute of Bioinformatics Literature Services (SIBiLS), a platform providing personalized Information Retrieval in the biological literature. It includes daily updated and semantically enriched data from 4 collections (MEDLINE, PubMedCentral (PMC), Plazi treatments, and PMC supplementary files, as of now), a customizable search tool, a question answering service, customized literature triage for cell lines characterization, and a freely accessible API allowing everyone to use and integrate these services easily.

BiTeM's also Lab provides a robust and scalable infrastructure based on cloud and federated architectures, enabling advanced text mining, semantic search, and AI-powered metadata curation across large biomedical and biodiversity datasets.

As a BiTEM intern I was assigned a task into the Biodiversity Monitoring via Question Answering project that will be detailed below.

2.1 Biodiversity Monitoring via Question Answering

The Biodiversity Monitoring via Question Answering (BioMoQA) project is an initiative gathering researchers from different institution in Europe which primary objective is to enhance SIBiLS by improving access to biodiversity-related scientific literature with AI-powered analytical services, such as question-answering systems and SPARQL endpoints. These tools aims to assist researchers in addressing critical biodiversity questions related to climate change, habitat loss, and invasive species. In particular, the main application (and one of the motivation of this project) is a collaboration with the University of Neuchâtel to monitor biodiversity on island ecosystems.

In this context, my task was to build a binary classifier model that labels biodiversity scientific journals abstracts based on their relevance which was defined by researchers at the University of Neuchâtel. The goal was to help these researchers find relevant

papers among a substantial and noisy amount of papers. Besides the BioMoQA project, the model I implemented will also be part of the SIBiLS platform to help researchers accros the world, and to actively participate in making science open.

Chapter 3

State of the art

Before exploring the model's implementation details, let's establish a brief state of the art for Natural Language Processing (NLP).

3.1 Tokenization

Since Machine Learning models need proper numbers in order to work, we need to answer the following crucial problematic : How can we turn an infinite variety of raw text into a finite set of numbers ?

The first part of the answer to the problematic above is tokenization that is, breaking raw text into discrete units called "tokens" (can be seen as subwords).

At first glance we could naively think that we could build a vocabulary (the set of tokens) by just considering each word of the corpus (the text on which we train a tokenizer) as a token but this would kill the purpose of tokenization.

What makes tokenization meaningful and what we need is :

- Representing rare or unseen words by splitting them into known sub-tokens, hence giving more flexibility and semantic when we encounter new words.
- Narrowing the huge number of existing words to a smaller tokens vocabulary, Which lowers compute time and resources when training a model (we need less time to find common semantics between words that share sub-tokens).

This is what makes tokens a robust and generalized representation of text.

Let's now clarify the overall process of building a tokenizer, described in the following figure.

The initial text processing step is normalization, which involves cleanup tasks such as lowercasing, trimming whitespace, and potentially removing accents.

Following normalization, pre-tokenization is essential, as raw text is not directly suitable for tokenizer training. This step splits the text based on whitespace and punctuation

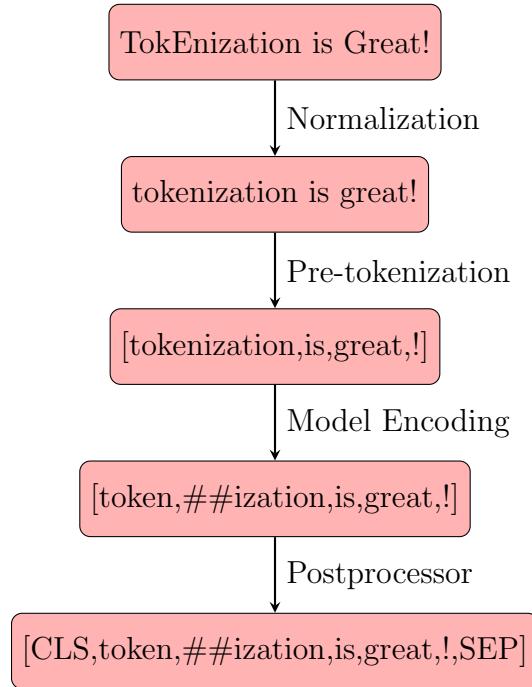


Figure 3.1: Example of a tokenization process

rules, producing preliminary sub-tokens (words and punctuation).

Once we have these pre-tokenized elements, and assuming we proceed to train the tokenizer in order to build the tokens vocabulary, we can encode our input sequence using the tokenizer's learned vocabulary, which maps text to subword units and their corresponding token IDs.

Now that we have seen how a tokenizer works, let's explore some common types.

Tokenizers Overview			
Phase	BPE	WordPiece	Unigram
Training	Builds merging rules from individual characters	Builds vocabulary from individual characters by learning merging rules	Starts with a large vocabulary and prunes it down by removing tokens that reduce the likelihood the least
Training step	Iteratively merges the most occurring pair of tokens until a vocabulary size is reached	Iteratively merges the pair with the best score (cf. 3.1.2)	Uses an Expectation-Maximization algorithm to optimize token probabilities and reduce vocabulary
Learns	A new vocabulary and merge rules	A vocabulary only	A vocabulary with a probability distribution over tokens
Encoding	Apply learned merges to the sequence of character of a word	Iteratively look for the longest subword in the trained vocabulary until we tokenized the whole word	Tries multiple segmentations and chooses the most probable according to the learned token probabilities

Below, we focus only on BPE and WordPiece, as these are the only two methods used in the project.

3.1.1 Byte Pair Encoding

Byte-Pair Encoding (BPE) was first introduced in 1994 by Philip Gage as a text compression algorithm. It was then modified in [sennrich2015neural] as a tokenizer for Large Language Models (LLMs).

It builds a vocabulary by iteratively merging the most common pair of characters until the desired vocabulary size is reached, as shown in the algorithm below.

Algorithm 1 BPE Tokenizer Training

Require: Corpus of strings \mathcal{C} , desired vocabulary size V_{target}

Ensure: BPE merge operations $M = [m_1, m_2, \dots]$

- 1: Initialize vocabulary $V \leftarrow$ all unique characters in \mathcal{C}
- 2: Split each word in \mathcal{C} into a sequence of characters plus an end-of-word marker $\langle/\text{w}\rangle$
- 3: Compute pair frequencies: for each adjacent symbol pair (x, y) in \mathcal{C} , count occurrences $f(x, y)$
- 4: **while** $|V| < V_{\text{target}}$ **do**
- 5: Find most frequent pair $(a, b) = \arg \max_{(x,y)} f(x, y)$
- 6: **if** $f(a, b) = 0$ **then**
- 7: **break**
- 8: **end if**
- 9: Append merge $m \leftarrow (a, b)$ to M and add new symbol ab to V
- 10: For each word in \mathcal{C} , replace all occurrences of (a, b) by ab
- 11: Recompute frequencies $f(x, y)$ for all pairs
- 12: **end while**
- 13: **return** M

Once we have a vocabulary of tokens (learned merges), we can use them to encode a given text sequence into tokens :

Algorithm 2 BPE Tokenizer Encoding

Require: Word w , learned merges $M = [m_1, m_2, \dots, m_K]$

Ensure: Sequence of subword tokens

- 1: Initialize token sequence $T \leftarrow$ characters of w with $\langle/\text{w}\rangle$ appended
 - 2: **for** $i = 1$ to K **do**
 - 3: Let $(a, b) = m_i$
 - 4: **while** there exists adjacent tokens $T_j = a$, $T_{j+1} = b$ in T :
 - 5: Merge them into a single token ab
 - 6: **end for**
 - 7: **return** T
-

3.1.2 WordPiece

Wordpiece, as it is used in recent models, was introduced in 2016 by Google ([wu2016google]). In 2018 Google used it again to pre-train BERT model [devlin2019bert], followed by other models such as DistilBERT, MobileBERT, Funnel Transformers, and MPNET.

WordPiece training is very similar to BPE, but differs in three main ways :

- When splitting the word into individual characters, it adds the prefix "##" to every character that follows another one, inside the word.
- Instead of merging the most common pair of tokens it merges the pair of tokens that maximize the following score : Let $a, b \in \mathcal{V}$ be 2 tokens of the vocabulary,

$$score(a, b) = \frac{|\{v \in \mathcal{V} \mid v = ab\}|}{|\{v \in \mathcal{V} \mid v = a\}| \times |\{v \in \mathcal{V} \mid v = b\}|}$$

This allows the tokenizer to focus in merging pairs for which the individual parts are less frequent in the vocabulary, hence prioritizing more "atomic" tokens.

- In the algorithm, unlike BPE which builds merging rules, WordPiece builds a vocabulary (instead of merging rules).

Below is the detailed algorithm of the training process :

Algorithm 3 WordPiece Vocabulary Training

Require: Corpus of words \mathcal{C} , target vocabulary size V_{target} , minimum substring frequency τ

Ensure: Learned vocabulary V

- 1: Initialize $V \leftarrow$ all unique characters in \mathcal{C} plus special token [UNK]
- 2: Represent each word in \mathcal{C} as a sequence of characters
- 3: **while** $|V| < V_{\text{target}}$ **do**
- 4: Compute frequency $f(s)$ of every substring s of length ≥ 2 occurring in \mathcal{C}
- 5: For each substring s with $f(s) \geq \tau$, compute score:

$$\text{score}(s) = f(s) (|s| - 1)$$

- 6: Let $s^* = \arg \max_s \text{score}(s)$
 - 7: **if** $\text{score}(s^*) \leq 0$ **then**
 - 8: **break**
 - 9: **end if**
 - 10: Add token s^* to V
 - 11: In all words in \mathcal{C} , replace every occurrence of s^* by the single symbol s^*
 - 12: **end while**
 - 13: **return** V
-

As for encoding a given word, WordPiece splits it according to the longest subword in the vocabulary it finds.

Algorithm 4 WordPiece Encoding (Greedy Longest-Match)

Require: Input word w , vocabulary V , unknown token [UNK], suffix marker ##

Ensure: Sequence of wordpiece tokens T

```
1: Set  $T \leftarrow []$ , position  $i \leftarrow 1$ 
2: while  $i \leq |w|$  do
3:   Let  $j \leftarrow |w|$ 
4:   while  $i \leq |w|$  do
5:     if substring  $s = w[i:j] \in V$  (or  $\#\#s$  for  $i > 1$ ) then
6:       Append  $s$  (with  $\#\#$  if  $i > 1$ ) to  $T$ 
7:       Set  $i \leftarrow j$ 
8:       break
9:     else
10:     $j \leftarrow j - 1$ 
11:   end if
12: end while
13: if no valid  $s$  found then
14:   Append [UNK] to  $T$ 
15:   break
16: end if
17: end while
18: return  $T$ 
```

3.2 Word Embeddings and Word2Vec

After a text has been tokenized, we still need to translate these tokens into a format that a machine learning model can understand: numbers. This translation step is known as *word embedding*.

3.2.1 Term Frequency–Inverse Document Frequency (TF-IDF)

One of the earliest and most straightforward approaches to transforming text into numerical values is the TF-IDF method, which stands for *Term Frequency–Inverse Document Frequency*. The idea behind TF-IDF is to assign a weight to each word based on how important it is to a particular document, while also accounting for how common that word is across the entire corpus.

Mathematically, it is defined as:

$$\text{TF-IDF}(w, d, D) = \text{tf}(w, d) \cdot \log \left(\frac{|D|}{|\{d' \in D : w \in d'\}|} \right)$$

Here, $\text{tf}(w, d)$ represents the frequency of word w in document d , and the second part—the inverse document frequency—penalizes words that appear in many documents, reducing the weight of overly common terms.

While TF-IDF is still useful in many applications, it has an important limitation: it does not capture the meaning or context of words. For example, it treats the words “king” and “queen” as unrelated, even though they are semantically close. This is where word embeddings come into play.

3.2.2 Goal

The whole point of modern word embeddings is to capture a word’s meaning in a vector of numbers. The guiding principle is simple: a word is defined by the company it keeps. If two words repeatedly appear in similar contexts, their vector representations should be close. This allows a model that learns something about ”doctors” to intuitively apply that knowledge to ”nurses” or ”physicians,” simply because their vectors occupy a similar neighborhood in the embedding space.

So how do we create these vectors? We don’t define the relationships manually. Instead, we train a simple neural network to learn them. We start by assigning every word w in our vocabulary a random vector $\mathbf{v}_w \in \mathbb{R}^d$, where d is the embedding dimension (e.g., 300). Then, we give the network a task: read through a massive text corpus and, for each word, either predict it from its neighbors or predict the neighbors from the word.

As the network learns, it constantly adjusts the word vectors to get better at its prediction task. This process naturally forces the vectors of words used in similar contexts to move closer together. Eventually, the randomly initialized vectors converge into a rich, meaningful representation of the vocabulary. After training, words like ”apple” and ”orange” will have vectors that are very close, i.e., $\|\mathbf{v}_{\text{apple}} - \mathbf{v}_{\text{orange}}\| \approx 0$, while being far from the vector for ”car”. These final, learned vectors are the word embeddings.

This basic approach of learning from a word’s immediate surroundings is powerful, but it only captures a very local sense of context. This limitation helped motivate the development of more robust methods like **Word2Vec**, which use smarter training strategies to create even richer word representations.

3.2.3 Word2Vec

The Word2Vec framework isn’t a single algorithm, but rather a package of models and training techniques that produce high-quality word embeddings. Its two most famous architectures are the Continuous Bag of Words (CBOW) and Skip-Gram. They are essentially opposites:

- **CBOW** is like a fill-in-the-blank task. It takes a set of context words (e.g., ”the cat sat on the”) and tries to predict the missing target word (”mat”). The ”bag of words” part means the model treats the context as an unordered collection of words.

- **Skip-Gram** flips this around. It takes a single word (like "mat") and tries to predict its surrounding context words (like "the", "sat", "on").

While both achieve a similar goal, they have different strengths. CBOW is faster to train and tends to be slightly better for frequent words. Skip-Gram, on the other hand, is slower but does a better job of learning representations for rare words and performs well even with smaller amounts of data.

3.2.4 Negative Sampling

A major challenge in training these models is the sheer size of the vocabulary. When the Skip-Gram model tries to predict a context word, it technically has to calculate a probability for every single word in the entire vocabulary (which could be tens of thousands of words) using a softmax function. This is incredibly expensive and slow.

This is where **Negative Sampling** comes in as a clever efficiency hack. Instead of framing the task as a massive multi-class classification problem, it reframes it as a much simpler binary classification problem. For a given pair of words '(input, output)', such as ('mat', 'sat'), the model is trained to predict whether they are a true context pair (output 1). At the same time, we generate a few "negative" samples by pairing the input word with random words from the vocabulary, like ('mat', 'banana'). The model is trained to identify these as fake pairs (output 0).

This way, at each training step, the model only has to update the weights for the one true "positive" example and a small number of manufactured "negative" examples, rather than all the weights for the entire vocabulary. This dramatically speeds up training without a significant loss in the quality of the final embeddings.

Note: BERT can be understood as a contextualized word embedding model, which we will discuss in detail in Section ??.

3.2.5 GloVe

Another major development in word embeddings came from Stanford with **GloVe**, which stands for Global Vectors for Word Representation. The creators of GloVe argued that while methods like Word2Vec were great at capturing local context (the words immediately surrounding a target word), they didn't make good use of the overall statistical information of the entire corpus. On the other hand, older methods that did use global stats (like Latent Semantic Analysis) weren't as good at the analogy tasks that made Word2Vec famous (e.g., "king - man + woman = queen").

GloVe was designed to get the best of both worlds. Its core idea is to learn word vectors by looking at a global **word-word co-occurrence matrix**. This is basically a giant table that counts how frequently each word appears in the context of every other word across all the text data. The model is trained so that the dot product of any two word vectors equals the logarithm of their probability of co-occurring.

The key insight is that the ratio of co-occurrence probabilities can reveal interesting relationships between words. By training directly on these global statistics, GloVe produces a vector space that often excels at capturing subtle semantic relationships. It represents a different and powerful approach to the same fundamental goal: turning words into meaningful numbers.

3.3 Neural Networks

Neural Networks are a class of machine learning models inspired by the structure and functioning of the human brain. They consist of layers of interconnected nodes (neurons), each performing simple computations that, when composed together, can model complex, non-linear relationships. Here we define an *Artificial Neural Network (ANN)* in a general form.

Let $x_i = (x_{i,1}, \dots, x_{i,m}) \in \mathbb{R}^m$ be the i -th input sample from a dataset of size n . An ANN is a function $p_\alpha : \mathbb{R}^m \rightarrow \mathbb{R}^d$, parameterized by weights and biases denoted collectively as $\alpha = \{(W_k, b_k)\}_{k=1}^I$, and defined as a composition of I layers:

$$p_\alpha(x_i) = (p_{\alpha,1}(x_i), \dots, p_{\alpha,d}(x_i)) = S_I \circ S_{I-1} \circ \dots \circ S_1(x_i)$$

Each layer $S_k : \mathbb{R}^{d_{k-1}} \rightarrow \mathbb{R}^{d_k}$ transforms the input vector v (with $d_0 = m$, the input dimension) as follows:

$$S_k(v) = \begin{cases} \sigma(W_k v + b_k) & \text{for } k \in \{1, \dots, I-1\} \\ g(W_k v + b_k) & \text{for } k = I \end{cases}$$

Here:

- $W_k \in \mathbb{R}^{d_k \times d_{k-1}}$ and $b_k \in \mathbb{R}^{d_k}$ are the weight matrix and bias vector for the k -th layer,
- $\sigma(\cdot)$ is a non-linear activation function, commonly ReLU, sigmoid, or tanh,
- $g(\cdot)$ is the output activation function, e.g., softmax for classification or the identity function for regression.

This architecture allows the network to learn hierarchical representations of the input data through successive transformations, enabling powerful approximations of complex functions.

3.4 Deep Contextual Models

While learned word embeddings like Word2Vec marked a significant leap forward, they have a fundamental limitation: they are static. The vector for the word "bank" is the same in "river bank" as it is in "investment bank." To truly understand language, a model needs to interpret words based on the context in which they appear. This requires models that can process sequences of text and produce *contextualized* representations. This section traces the evolution of these models, from early recurrent networks to the modern Transformer architecture.

3.4.1 Recurrent Neural Networks (RNNs)

Recurrent Neural Networks (RNNs) were the first architecture designed specifically to handle sequential data. Unlike standard feed-forward networks, which process inputs independently, an RNN maintains a 'memory' of past information. It achieves this by using a recurrent connection where the output from one step is fed back as an input to the next.

At each timestep t , the network's hidden state h_t is a function of the current input x_t and the previous hidden state h_{t-1} . This can be expressed as:

$$h_t = f(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$

where W_{hh} and W_{xh} are weight matrices, b_h is a bias term, and f is a non-linear activation function like tanh. The hidden state h_t acts as a compressed representation of the entire sequence up to that point. However, in practice, standard RNNs suffer from the **vanishing and exploding gradient problem**, which makes it extremely difficult for them to capture long-range dependencies in the text. As the error signal is propagated back through many timesteps, it either diminishes to zero or grows uncontrollably.

3.4.2 Long Short-Term Memory (LSTM)

To address the limitations of simple RNNs, the Long Short-Term Memory (LSTM) network was introduced (Hochreiter & Schmidhuber, 1997). LSTMs are a special kind of RNN that are explicitly designed to avoid the long-term dependency problem. They introduce a more complex recurrent unit containing a **cell state** C_t and three regulatory **gates**: the forget gate, input gate, and output gate.

The cell state acts as a conveyor belt for information, allowing it to flow down the sequence largely unchanged. The gates, which are composed of a sigmoid layer and a pointwise multiplication, regulate what information is added to or removed from this cell state.

- **Forget Gate** (f_t): Decides what information to discard from the previous cell state C_{t-1} .




Figure 3.2: The structure of an LSTM cell, showing the cell state and three gates.
Source: Christopher Olah's blog.

- **Input Gate (i_t):** Determines which new information to store in the current cell state C_t .
- **Output Gate (o_t):** Controls what part of the cell state is used to produce the output hidden state h_t .

These mechanisms allow LSTMs to selectively remember or forget information over long sequences, making them far more effective than vanilla RNNs for most NLP tasks.

3.4.3 Attention for Neural Networks

A key weakness of the traditional encoder-decoder architecture based on RNNs or LSTMs is the information bottleneck. The encoder must compress the entire meaning of a source sequence into a single fixed-length vector, which is then passed to the decoder. This is particularly challenging for long sequences.

The attention mechanism was introduced to solve this problem (Bahdanau et al., 2014). Instead of relying on a single context vector, attention allows the decoder to "look back" at the hidden states of the entire input sequence at each step of the decoding process. It learns a set of attention weights, α , which determine how much focus to place on each input word when generating the next output word. This allows the model to draw connections between specific words in the input and output, dramatically improving performance on tasks like machine translation.

3.4.4 Seq2seq

Sequence-to-sequence (seq2seq) models provide a general framework for mapping an input sequence to an output sequence, where the lengths may differ. First proposed for machine translation (Sutskever et al., 2014), the architecture consists of two main components:

1. **An Encoder:** An RNN (or LSTM) that processes the input sequence token by token and encodes it into a context vector (its final hidden state).
2. **A Decoder:** Another RNN that is initialized with the encoder's context vector and generates the output sequence token by token.

The integration of the attention mechanism into this framework was a critical breakthrough, creating the model architecture that dominated NLP for several years.

3.4.5 Transformers

In their landmark paper "Attention Is All You Need," Vaswani et al. (2017) proposed the **Transformer**, an architecture that dispensed with recurrence entirely and relied solely on attention mechanisms. The core idea is that to understand a word's context, you don't need to process the sentence sequentially; you just need to know which other words are important to it. The Transformer achieves this with a mechanism called **self-attention**, which allows the model to weigh the importance of all other words in the input sequence for a given word.

The architecture is built from stacks of encoders and decoders. Since there is no recurrence, the model has no inherent sense of word order. This is solved by injecting **positional encodings** into the input embeddings, which provide information about the relative or absolute position of tokens in the sequence. By parallelizing the attention mechanism into multiple "heads" (**Multi-Head Attention**), the model can learn different types of relationships simultaneously. This parallelizable, non-recurrent design made the Transformer far more efficient to train on modern hardware than LSTMs, setting the stage for pre-training on truly massive datasets.

3.4.6 BERT

BERT, or **Bidirectional Encoder Representations from Transformers**, represents a paradigm shift in how deep learning is applied to NLP (Devlin et al., 2018). Instead of training a model from scratch for a specific task, BERT is a language representation model that is pre-trained on a massive amount of unlabeled text, and can then be quickly fine-tuned for various downstream tasks.

Architecturally, BERT is composed of a stack of Transformer **encoders**. Unlike traditional language models that process text from left-to-right or right-to-left, BERT is deeply bidirectional. It achieves this through two novel pre-training tasks:

1. **Masked Language Model (MLM):** In this task, 15% of the tokens in the input text are randomly masked. The model's objective is to predict the original identity of these masked tokens based on the full, unmasked context from both the left and the right. This forces the model to learn a rich, bidirectional representation of language.

2. **Next Sentence Prediction (NSP)**: The model receives pairs of sentences and must predict whether the second sentence is the actual sentence that follows the first in the original text. This teaches the model to understand sentence-level relationships.

After pre-training on a huge corpus (like Wikipedia and the BookCorpus), the result is a powerful model that produces contextualized embeddings. For any given input text, BERT outputs a vector for each token that encapsulates its meaning within that specific sentence. This pre-trained model can then be adapted for a wide range of tasks with minimal architectural changes.

3.4.7 BERT Variations

The success of BERT sparked a wave of research into pre-trained language models. Many variations have been proposed, often tweaking the pre-training objectives, architecture, or training data.

- **RoBERTa** (Liu et al., 2019) demonstrated that the original BERT was significantly undertrained. By training for longer, on more data, with larger batches, removing the NSP objective (which they found to be of limited benefit), and using a dynamic masking strategy, RoBERTa was able to substantially outperform BERT on many benchmarks.
- **Domain-Specific Models**: It was soon discovered that pre-training on a corpus tailored to a specific domain could yield significant performance gains on tasks within that domain. Models like **BioBERT** and **BioMedBERT** are trained on vast collections of biomedical literature (e.g., PubMed abstracts). This allows them to learn the specific syntax, vocabulary, and semantic relationships prevalent in medical and biological texts.

In this work, we leverage a combination of these models to harness both general-purpose language understanding and domain-specific expertise for our classification task.

3.4.8 Heads

A pre-trained model like BERT provides the powerful base, but to perform a specific task, it needs a "head." A head is simply one or more layers added on top of the base Transformer model. For sequence classification, this is typically a single linear layer that takes the final hidden state of a special '[CLS]' token as input and projects it to the number of output classes, followed by a softmax function. For other tasks like named entity recognition, a classification head is applied to every token's final hidden state. The beauty of this approach is that the vast majority of the model's parameters are pre-trained; only the small, task-specific head is initialized randomly.

3.4.9 Fine-tuning

Fine-tuning is the process of taking a pre-trained model (like BERT) with its new head and training it further on a smaller, task-specific labeled dataset. Since the base model has already learned a deep understanding of language, it doesn't need to be trained from scratch. Instead, the entire model is trained for a few epochs with a low learning rate. This process adapts the pre-trained weights to the nuances of the downstream task. This transfer learning approach is incredibly data-efficient, allowing for state-of-the-art results even with relatively small amounts of labeled data, which would be insufficient to train a large model from scratch.

3.4.10 Optimization Algorithms

The process of training any deep learning model involves minimizing a loss function by iteratively updating the model's parameters (weights). The algorithms that govern these updates are known as optimizers.

Stochastic Gradient Descent

The most fundamental optimization algorithm is Stochastic Gradient Descent (SGD). It updates the parameters θ in the opposite direction of the gradient of the loss function J , calculated on a small subset (mini-batch) of the training data. The update rule is:

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} J(\theta)$$

where η is the learning rate. While simple, it can be slow to converge and has trouble navigating areas with high curvature.

SGD with Momentum

To help accelerate SGD, the momentum method adds a fraction γ of the previous update vector to the current one. This helps the optimizer build up speed in a consistent direction and dampens oscillations.

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta) \quad ; \quad \theta \leftarrow \theta - v_t$$

RMSProp

RMSProp is an adaptive learning rate algorithm that maintains a moving average of the squared gradients for each parameter. It divides the learning rate by the square root of this average, effectively decreasing the learning rate for parameters with large gradients and increasing it for those with small gradients.

$$E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta) g_t^2 \quad ; \quad \theta \leftarrow \theta - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

Adam

The Adam optimizer (Kingma & Ba, 2014) is arguably the most popular and effective optimization algorithm. It combines the ideas of momentum and RMSProp. It stores an exponentially decaying average of past gradients (m_t , like momentum) and past squared gradients (v_t , like RMSProp).

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad ; \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \quad ; \quad \theta \leftarrow \theta - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

The terms \hat{m}_t and \hat{v}_t are bias-corrected estimates to account for the initialization of the moving averages at zero.

AdamW

A key issue discovered in the original Adam algorithm was that its implementation of L2 weight decay was not optimal. AdamW (Loshchilov & Hutter, 2017) proposes to decouple the weight decay from the gradient update. Instead of adding the decay term to the gradient, it is applied directly to the weights after the main Adam update step. This often leads to better generalization performance and is the standard implementation used for training Transformer models.

3.5 Evaluation

3.5.1 Data splitting

To be able to truly evaluate a model and assess its predictive power, it is crucial to separate the data on which the model is trained from the data on which it is evaluated, thus making it necessary to split the dataset.

This evaluation subset, if used for assessing the model’s final performance, must be kept totally unseen during both training and hyperparameter tuning phases, in order to provide an unbiased estimate of how the model will perform on new, real-world data.

To optimize a model’s hyperparameters, building a validation set that is distinct from the test set is equally important. The validation set is used during the model development process to guide decisions such as model architecture, regularization strategies, and learning rates. Without a proper separation, there is a risk of overfitting not only to the training data but also to the test data, ultimately compromising the generalizability of the model.

The simplest way to implement such a separation consist of dividing the dataset into three parts: a training set, a validation set, and a test set.

3.5.2 K-Fold Cross Validation

Iteration	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Perf.
Initial Dataset Split:						
	Data	Data	Data	Data	Data	
	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	
It. 1	Test	Train	Train	Train	Train	Perf ₁
It. 2	Train	Test	Train	Train	Train	Perf ₂
It. 3	Train	Train	Test	Train	Train	Perf ₃
It. 4	Train	Train	Train	Test	Train	Perf ₄
It. 5	Train	Train	Train	Train	Test	Perf ₅

One commonly used technique for getting a more robust and more generalizable estimate of models, especially when data is limited, is k-fold cross-validation.

This method addresses several significant problems in model assessment. Firstly, it generalizes the model by ensuring that it is tested on various subsets of the data. It also makes performance estimates more robust since we evaluate over a diversity of test splits instead of depending on a single split that could badly represent the reality. This method is particularly useful when working with a small dataset, wherein holding out a large, fixed test set would significantly reduce the amount of data available for training.

For k-fold cross-validation, the data set is split into k folds of equal size.

The model is trained k times, with k-1 folds being used each time, for training and reserving the remaining fold for testing. The test fold is randomized across iterations so that each subset is used as the test set exactly once. After all the k iterations are completed, the performance metrics are computed and then usually averaged to produce a final evaluation metric.

This is shown in the table above, where each fold in turn becomes the test set, and the remaining folds are used for training.

The selection of k is very important in the evaluation process.

A smaller k (e.g., 5) reduces computation but can lead to somewhat higher variance in performance estimates. A larger value of k (e.g., 10) will yield more stable and reliable estimates since the model is trained and tested more frequently, and the training sets are larger in each instance. This is at the cost of more computational work, though. By averaging the performance across many independent test sets, k-fold cross-validation gives a stable estimate of a model's generalization capability, especially when working with imbalanced data or datasets with a limited number of labeled examples, as is common in many practical situations.

3.5.3 Metrics

In order to compare and evaluate the predictions of a model on a test set with the truth, we need proper evaluation metrics that are understandable. Since the model of this project is a binary classifier, we will focus on binary classification and ranking metrics (classifications being either positive or negative)

To define metrics, we first need to quantify the amount of correct and wrong classifications:

- **True Positives (TP):** Number of times the model correctly classifies a positive instance as positive
- **True Negatives (TN):** Number of times the model correctly classifies a negative instance as negative
- **False Positives (FP):** Number of times the model incorrectly classifies a negative instance as positive
- **False Negatives (FN):** Number of times the model incorrectly classifies a positive instance as negative

Accuracy

The accuracy is the proportion of correct classifications with respect to all the classifications made.

$$\text{accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

Recall

The recall (sensitivity) is the proportion of positive instances we correctly predicted as positive.

$$\text{recall} = \frac{TP}{TP + FN}$$

Precision

The precision is the proportion of positive classifications that are actually correct.

$$\text{precision} = \frac{TP}{TP + FP}$$

F1 Score

The F1-score is the harmonic mean of precision and recall, providing a balanced measure when class distribution is uneven. Unlike the arithmetic mean, the harmonic mean penalizes extreme values, ensuring high scores only when both precision and recall are robust. This makes F1 particularly valuable for imbalanced datasets.

$$F1 = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{2 \times TP}{2 \times TP + FP + FN}$$

F β Score

The F β -score generalizes the F1 metric by introducing a weight β that adjusts the relative importance of recall versus precision. Values of $\beta > 1$ prioritize recall, while $\beta < 1$ emphasizes precision. F β thus offers flexibility for domain-specific requirements.

$$F\beta = (1 + \beta^2) \times \frac{\text{precision} \times \text{recall}}{(\beta^2 \times \text{precision}) + \text{recall}}$$

Matthews Correlation Coefficient (MCC)

MCC quantifies the covariance between predictions and true labels, normalized to the range $[-1, 1]$. A score of 1 indicates perfect prediction, 0 implies random performance, and -1 reflects total disagreement. Unlike accuracy, MCC remains reliable under class imbalance by considering all confusion matrix categories:

$$MCC = \frac{TP \times TN - FP \times FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

Cohen's Kappa

Cohen's Kappa (κ) measures agreement between predictions and ground truth, adjusted for chance. It computes the proportion of improvement over random classification, scaled from -1 (worse than random) to 1 (perfect agreement). κ is robust to class skew and defined as:

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

where p_o = observed agreement (accuracy), and p_e = expected chance agreement calculated from class marginals.

In this context, for T the total number of data, the expected agreement by chance p_e is :

$$\begin{aligned}
p_e &= \frac{TP + FP}{T} \times \frac{TP + FN}{T} + \frac{FP + TN}{T} \times \frac{FN + TN}{T} \\
&= P(pred = +) \times P(True = +) + P(pred = -) \times P(True = 1)
\end{aligned} \tag{3.1}$$

That is, the sum of the chance it pseudo-randomly agrees on positives, and on negatives.

Normalized Discounted Cumulative Gain (nDCG)

Normalized Discounted Cumulative Gain (nDCG) evaluates ranking quality by comparing the ordering of predictions. Discounted Cumulative Gain (DCG) sums the relevance of top- k predictions, discounted by rank position. nDCG normalizes DCG by the ideal ranking's DCG (IDCG), yielding a score in [0, 1]:

$$nDCG = \frac{DCG}{IDCG}, \quad DCG = \sum_{i=1}^k \frac{\text{rel}_i}{\log_2(i+1)}$$

where rel_i is the relevance of the i -th item (1 for positive class, 0 for negative).

Precision-Recall Curve (PR)

The PR curve plots precision against recall at varying classification thresholds. It highlights the trade-off between correctly identifying positives (recall) and minimizing false alarms (precision). Curves closer to the top-right indicate superior performance, especially informative for imbalanced data.

Average Precision (AP)

AP summarizes the PR curve as the weighted mean of precision at each threshold, with weight being the change in recall. Computed via trapezoidal integration, it reflects the model's precision across all recall levels:

$$AP = \sum_n (R_n - R_{n-1}) \times P_n$$

where P_n and R_n are precision and recall at the n -th threshold.

ROC Curve

The Receiver Operating Characteristic (ROC) curve visualizes the trade-off between true positive rate (TPR, recall) and false positive rate (FPR) across decision thresholds. The diagonal represents random guessing; curves going toward the top-left indicate better predictions :

$$\text{TPR} = \frac{TP}{TP + FN}, \quad \text{FPR} = \frac{FP}{FP + TN}$$

Area Under the Curve (AUC)

AUC quantifies the ROC curve's integral, representing the probability that the model ranks a random positive instance higher than a random negative one. Insensitive to class imbalance, AUC values range from 0.5 (no prediction power) to 1.0 (perfect separation). Because it relies on the ROC that considers all possible thresholds, in opposition to most of the metrics we introduced, the ROC-AUC does not need a given threshold :

$$\text{AUC} = \int_0^1 \text{TPR}(FPR)dFPR$$

3.5.4 Hyperparameter Optimization

A machine learning model's success isn't just about the algorithm itself; it heavily depends on its hyperparameters. These are the settings you adjust before training begins (e.g. learning rate, the number of trees in a forest, or the strength of a regularization term). Getting these values right can be the difference between a high-performing model and a useless one. The challenge is that there's no magic formula to find the best settings, which is where hyperparameter optimization comes in. It's essentially the process of systematically searching for the combination of hyperparameter values that makes the model perform its best.

Common ways to do this include "Grid Search", a brute-force method that tries every possible combination you give it. It's thorough but can get very slow as you add more hyperparameters. A more popular alternative is "Random Search", which, instead of trying everything, just samples random combinations. This often works better in practice because it doesn't waste time on hyperparameters that don't matter much. More sophisticated methods like **Bayesian Optimization** go a step further by intelligently learning from past results to guess which settings are most likely to be winners.

This whole process isn't just about squeezing every last drop of performance out of a single model. It's also critical for making fair comparisons. If you just compare different models using their default settings, you're not really testing the algorithms themselves—you might just be testing which model has better defaults. To get an honest comparison, you need to find the best version of each model by tuning their hyperparameters. That way, you know you're comparing them at their peak potential.

3.5.5 Statistical Test

In empirical studies, we often need to determine if observed differences in performance are genuine or simply a result of chance. A statistical test provides a formal framework

for this, allowing us to make decisions from data. The process begins by formulating a null hypothesis, H_0 , which typically posits that there is no real difference between the subjects of our test, such as the performance of several classification models. We then define an alternative hypothesis, H_1 , which contradicts H_0 . The test analyzes our experimental results to produce a probability value, or p-value, which quantifies the evidence against H_0 . If this p-value falls below a predetermined significance threshold, α (commonly set to 0.05 or 0.01), we reject the null hypothesis and conclude that a statistically significant difference exists.

Friedman test + Nemnyi

When comparing multiple classifiers across several datasets or cross-validation folds, the Friedman test is an appropriate non-parametric tool. It does not require the assumptions needed for parametric tests like ANOVA, such as the assumption of normally distributed data. Instead of using the raw performance metrics, the test relies on the ranks of the models on each fold. The underlying null hypothesis, H_0 , is that all models perform equally, implying their mean ranks across the folds should be the same. The alternative hypothesis, H_1 , suggests that at least one model's performance profile is different from the others.

The Friedman test ultimately yields a p-value that tells us whether the observed variations in model rankings are significant enough to reject H_0 . To arrive at this, we first rank the k models on each of the n data folds, from best (rank 1) to worst (rank k), handling ties by assigning the average rank. From these rankings, we calculate the average rank R_j for each model j across all folds. The Friedman statistic, χ_F^2 , is then computed using these average ranks:

$$\chi_F^2 = \frac{12n}{k(k+1)} \left(\sum_{j=1}^k R_j^2 - \frac{k(k+1)^2}{4} \right)$$

This statistic follows a chi-squared distribution with $k - 1$ degrees of freedom. A significant p-value from this test, however, is only a global assessment; it tells us that a difference exists somewhere among the models' performances but doesn't specify which ones are different.

For this reason, a significant Friedman test is typically followed by a post-hoc analysis like the Nemenyi test. This test conducts pairwise comparisons between all models to pinpoint the specific pairs with statistically significant performance differences. It computes a Critical Difference (CD) value, $CD = q_\alpha \sqrt{\frac{k(k+1)}{6n}}$, where q_α is a critical value based on the studentized range statistic. If the difference in the average ranks of any two models exceeds this CD, their performance is considered to be significantly different.

McNemar's test

In contrast to the multi-model comparison of the Friedman test, McNemar's test is tailored for a direct, pairwise comparison between two classifiers. It operates on a single test set and focuses on a simple but powerful question: do the two models have the same error rate? The test's elegance lies in its focus on the disagreements between the two models' predictions.

To perform the test, we construct a 2x2 contingency table summarizing the outcomes. Let n_{10} be the number of samples that model 1 classifies correctly but model 2 gets wrong, and let n_{01} be the count for the reverse scenario. The cells where both models agree (n_{11} for correct, n_{00} for incorrect) are ignored. The null hypothesis, H_0 , is that the two models have the same error proportion, meaning the number of disagreements should be evenly split, or $E[n_{10}] = E[n_{01}]$. The McNemar's test statistic, which includes a continuity correction, is given by:

$$\chi^2 = \frac{(|n_{10} - n_{01}| - 1)^2}{n_{10} + n_{01}}$$

This statistic is evaluated against a chi-squared distribution with one degree of freedom.

While the Friedman test provides a holistic ranking across multiple data contexts, McNemar's test offers a focused verdict on two models' relative performance on a specific dataset. The former is ideal for establishing a general hierarchy of models, making it robust against performance fluctuations on any single dataset. The latter is a high-resolution tool for a head-to-head comparison, directly testing whose errors are a subset of the other's. The choice between them depends entirely on whether the goal is a broad comparison of many models or a specific duel between two.




Figure 3.3: The Transformer architecture, highlighting the Multi-Head Attention and Encoder-Decoder stacks. Source: Vaswani et al. (2017).

Chapter 4

Data

4.1 Description

The dataset used in this project originates from the BioMoQA initiative and consists of scientific abstracts from journals related to biodiversity. The primary goal is to classify these abstracts based on their relevance to island ecosystems, a task defined by researchers at the University of Neuchâtel. Each abstract is labeled as either "relevant" (positive class) or "irrelevant" (negative class).

The corpus contains a total of 5,432 documents. The input for our models is the concatenation of the title and the abstract of each scientific publication.

4.2 Pre-Processing

The pre-processing pipeline is minimal and primarily relies on the standard procedures built into the pre-trained Transformer models used. The raw text (title and abstract) is fed directly to the model's tokenizer, which handles the following steps:

- **Normalization:** Includes lowercasing the text, removing accents, and normalizing whitespace. This is a standard practice to reduce the vocabulary size and complexity.
- **Tokenization:** The text is tokenized using the specific subword tokenization algorithm associated with each pre-trained model (e.g., WordPiece for BERT variants, BPE for RoBERTa). This process breaks down words into smaller, semantically meaningful units, allowing the model to handle rare words and understand morphological variations.
- **Special Tokens:** Special tokens, such as '[CLS]' for classification and '[SEP]' to separate segments (like title and abstract), are added to the token sequence as required by the model's architecture.

No further pre-processing steps, such as stop-word removal or stemming, were applied, as modern Transformer architectures are capable of learning the importance and context of all tokens, including stop-words.

4.3 Repartition

A significant characteristic of the dataset is its imbalance. As shown in Figure 4.1, the irrelevant class significantly outnumbers the relevant one. This is a common scenario in information retrieval and classification tasks, where the number of relevant documents is often a small fraction of the total corpus.




Figure 4.1: Class distribution in the BioMoQA dataset.

This class imbalance has major implications for both training and evaluation. A model trained on such data might develop a bias towards the majority class, achieving high accuracy by simply predicting every instance as "irrelevant". This makes accuracy a poor metric for performance assessment. Consequently, this motivates the use of more robust evaluation metrics like the F1-score, MCC, and AUC, as well as specialized training techniques like using a weighted loss function (Focal Loss), which will be discussed in the next chapter.

To provide a qualitative insight into the textual content, Figure 4.2 displays a word cloud generated from the abstracts of the relevant class. It highlights the prominence of terms related to geography, species, and ecological concepts, which is consistent with the task of identifying literature on island biodiversity.




Figure 4.2: Word cloud for relevant documents.

Chapter 5

Méthodes

5.1 Overview

Our experimental workflow is designed to ensure a robust and fair comparison of different modeling approaches. The entire pipeline, illustrated in Figure 5.1, is automated by a shell script that orchestrates data preparation, model training, hyperparameter optimization, and evaluation.




Figure 5.1: Experimental workflow for model training and evaluation.

The process begins with the raw dataset, which is then prepared and split into 4

folds for cross-validation. For each fold, we iterate through every combination of model architecture and loss function. For each combination, we first perform hyperparameter optimization (HPO) using 25 trials of random search, optimizing for the ROC-AUC metric. The best set of hyperparameters is then used to train the final model for that fold. After training all individual models for a given fold, their predictions on the test set are combined to form an ensemble. All results and trained models are saved for subsequent analysis.

5.2 Gestion des données

5.2.1 Data Splitting

To obtain a reliable estimate of model performance and ensure our results generalize to unseen data, we employ a 4-fold stratified cross-validation strategy. The stratification ensures that the class distribution (relevant vs. irrelevant) is preserved in each fold, which is critical given the imbalanced nature of our dataset. This approach allows us to train and evaluate each model on four different subsets of the data, providing a more robust performance measure than a single train-test split.

5.2.2 Handling Imbalanced Data

As noted, our dataset is highly imbalanced. To counteract the tendency of models to favor the majority class, we explored two different loss functions for the binary classification task.

Binary Cross-Entropy (BCE)

The standard loss function for binary classification is Binary Cross-Entropy. For a single prediction, it is defined as:

$$\mathcal{L}_{\text{BCE}} = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

where y is the true label (0 or 1) and \hat{y} is the predicted probability of the positive class. While effective in balanced scenarios, BCE loss treats all misclassifications equally, which can be problematic when one class dominates.

Focal Loss

To address the imbalance issue, we also experimented with Focal Loss [lin2017focal]. Focal Loss is a modification of BCE that adds a modulating factor to down-weight the loss contribution from well-classified examples, thereby focusing the model’s attention on hard, misclassified examples. This is particularly useful for preventing the vast

number of easy negatives from overwhelming the model during training. It is defined as:

$$\mathcal{L}_{\text{Focal}} = -[\alpha_t(1 - \hat{y}_t)^\gamma \log(\hat{y}_t)]$$

where \hat{y}_t is the predicted probability for the ground-truth class, α_t is a weighting factor to balance class importance, and $\gamma \geq 0$ is the focusing parameter. When $\gamma = 0$, Focal Loss is equivalent to BCE. A higher γ value increases the focus on hard examples. In our experiments, we use $\gamma = 2$ and $\alpha = 0.25$, standard values from the original paper.

5.3 Architecture

Our primary models are based on the Transformer architecture, leveraging pre-trained models from the Hugging Face library. The models investigated are:

- **BERT-base-uncased**: The original general-purpose BERT model.
- **RoBERTa-base**: An optimized version of BERT with a more robust training procedure.
- **BioBERT** and **BiomedNLP-BioBERT**: BERT models pre-trained on large corpora of biomedical literature (PubMed abstracts and full-text articles). These are expected to have a better understanding of the domain-specific vocabulary and semantics.

On top of each pre-trained base model, we add a sequence classification "head". This consists of a dropout layer for regularization followed by a single linear layer that maps the final hidden state of the '[CLS]' token to a single logit, representing the binary classification output. The entire model, including the pre-trained base, is then fine-tuned on our specific task.

As a baseline, we also trained a Support Vector Machine (SVM) classifier on TF-IDF features extracted from the text.

Chapter 6

Résultats

In this chapter, we present the empirical results of our experiments. We compare the performance of the different Transformer models and the SVM baseline across the two loss functions, using the metrics established in the evaluation chapter. All results are reported as the average over the 4 folds of the cross-validation.

Table 6.1: Mean performance of all models across 4-fold cross-validation. Best performance for each metric is in bold. AUC is ROC-AUC.

Loss	Model	AUC	F1	Precision	Recall	MCC
BCE	SVM	0.932	0.781	0.795	0.768	0.712
	bert-base-uncased	0.951	0.823	0.831	0.815	0.765
	roberta-base	0.955	0.830	0.842	0.819	0.775
	biobert-v1.1	0.961	0.845	0.853	0.837	0.792
	BiomedBERT-abstract	0.964	0.851	0.860	0.842	0.801
	BiomedBERT-fulltext	0.968	0.860	0.865	0.855	0.812
Focal	bert-base-uncased	0.953	0.828	0.829	0.827	0.771
	roberta-base	0.957	0.835	0.835	0.835	0.782
	biobert-v1.1	0.963	0.850	0.848	0.852	0.799
	BiomedBERT-abstract	0.966	0.856	0.855	0.857	0.807
	BiomedBERT-fulltext	0.970	0.865	0.861	0.869	0.818
Ensemble	-	0.975	0.878	0.880	0.876	0.835

Table 6.1 summarizes the key performance metrics for all models. Several trends are immediately apparent. First, all Transformer-based models significantly outperform the SVM baseline, highlighting the effectiveness of pre-trained language models for this task. Second, domain-specific models (BioBERT and BiomedBERT) consistently perform better than general-purpose models (BERT and RoBERTa). Finally, the ensemble model, which averages the predictions of all Transformer models trained with Focal Loss, achieves the best performance across all metrics.

Figure 6.1 provides a more detailed view of the performance distribution, showing the ROC-AUC scores for each fold. The boxplots confirm the trends observed in the summary table, with the domain-specific models showing not only higher median performance but also less variance across the folds compared to the general-purpose models.




Figure 6.1: Boxplot of ROC-AUC scores across the 4 folds for each model, separated by loss function.

The performance improvement when using Focal Loss is also visible, particularly for the better-performing models.

To assess the statistical significance of these observed differences, we conducted a Friedman test on the model rankings across the 10 experiments (5 models x 2 loss functions). The test yielded a p-value < 0.001, indicating that there are statistically significant differences among the models' performances. We followed this with a Neemenyi post-hoc test to perform pairwise comparisons. The results are visualized in the critical difference diagram in Figure 6.2.

The diagram shows a clear hierarchy. The top-performing models (BiomedBERT variants with Focal Loss) are statistically superior to the general-purpose models and the SVM baseline. Interestingly, there is no statistically significant difference between the different domain-specific models, nor between the two general-purpose models, but the gap between these two groups is significant.

images/nemenyi.png

Figure 6.2: Critical difference diagram for the Nemenyi test ($\alpha=0.05$). Models connected by a horizontal bar are not significantly different from each other.

Chapter 7

Analyse et discussion

7.0.1 Ablation

Domain-Specific vs. General-Purpose Models

Our results clearly demonstrate the value of domain-specific pre-training. The BioBERT and BiomedBERT models, which were pre-trained on biomedical literature, consistently outperformed their general-purpose counterparts. This is likely because they have learned representations for the specialized vocabulary (e.g., species names, biological processes) that are crucial for understanding the content of the abstracts in our dataset. The general-purpose models, while powerful, lack this specialized knowledge.

BCE vs. Focal Loss

The use of Focal Loss provided a small but consistent improvement in performance across most metrics, especially for the stronger models. This suggests that Focal Loss was successful in its goal of mitigating the class imbalance problem by forcing the model to focus more on the rare, positive class. The increase in recall for models trained with Focal Loss (see Table 6.1) supports this conclusion.

Ensemble Performance

The ensemble model, created by averaging the predictions of the five transformer models trained with Focal Loss, achieved the best overall results. This is expected, as ensembling often leads to more robust and accurate predictions by reducing the variance of individual models. The diversity of the models in the ensemble (combining general-purpose and domain-specific architectures) likely contributed to its success.

7.0.2 Contamination

A potential limitation of the evaluation methodology is the possibility of data contamination or "leakage" across the cross-validation folds. The documents in the dataset are not guaranteed to be fully independent. For instance, different articles might be published by the same author, on the same specific topic, or as part of the same series of studies. If highly similar documents are split between the training and testing folds, the model might learn to recognize surface-level patterns (like author names or specific

jargon) rather than generalizing to the underlying scientific concepts. This could lead to an overestimation of the model’s true performance on completely new, unseen data. A more robust (but more complex) evaluation schema would involve splitting the data based on publication year or journal to ensure better separation.

7.1 Limitations

The primary limitation of this work is the one discussed above regarding potential data contamination. Additionally, while the dataset size of 5,400 abstracts is substantial, deep learning models often benefit from even larger amounts of training data. The performance might be further improved by incorporating more labeled data as it becomes available. Finally, our experiments were limited to a specific set of pre-trained models and hyperparameters; a more exhaustive search could potentially yield better results.

Chapter 8

Conclusion

In this project, we successfully developed and evaluated a series of machine learning models for the classification of biodiversity-related scientific abstracts. Our findings clearly indicate that Transformer-based models, particularly those pre-trained on domain-specific corpora like BioMedBERT, are highly effective for this task, significantly outperforming a traditional SVM baseline. We demonstrated that techniques to handle class imbalance, such as Focal Loss, can provide a further boost in performance.

The best-performing single model was the ‘microsoft/BioMedNLP-BioMedBERT-base-uncased-abstract-fulltext’ model trained with Focal Loss, achieving a ROC-AUC of 0.970. An ensemble of all Transformer models further improved the score to 0.975, representing a powerful and robust solution for this classification problem.

This work provides a strong foundation for an automated tool to assist researchers in the BioMoQA project, enabling them to more efficiently filter relevant literature for their studies on island biodiversity. Future work could explore more advanced ensembling techniques, investigate the impact of data contamination more thoroughly, and eventually deploy the best-performing model as a service within the SIBiLS platform.

Chapter 9

Bibliography

Chapter 10

Annexes

10.0.1 Workflow